

Importing college.txt and setup

```
college = read.table("college.txt")
college$Elite=as.factor(college$Elite)
college$Private=as.factor(college$Private)
attach(college)
summary(college)
```

```
## Private      Apps      Accept      Enroll      Top10perc
## No :212      Min.   :   81      Min.   :   72      Min.   :   35      Min.   : 1.00
## Yes:565      1st Qu.:  776      1st Qu.:  604      1st Qu.:  242      1st Qu.:15.00
##           Median : 1558      Median : 1110      Median :  434      Median :23.00
##           Mean   : 3002      Mean   : 2019      Mean   :  780      Mean   :27.56
##           3rd Qu.: 3624      3rd Qu.: 2424      3rd Qu.:  902      3rd Qu.:35.00
##           Max.   :48094      Max.   :26330      Max.   :6392      Max.   :96.00
## Top25perc    F.Undergrad    P.Undergrad      Outstate
## Min.   : 9.0      Min.   : 139      Min.   : 1.0      Min.   : 2340
## 1st Qu.:41.0      1st Qu.: 992      1st Qu.: 95.0      1st Qu.: 7320
## Median :54.0      Median :1707      Median : 353.0      Median : 9990
## Mean   :55.8      Mean   :3700      Mean   : 855.3      Mean   :10441
## 3rd Qu.:69.0      3rd Qu.:4005      3rd Qu.: 967.0      3rd Qu.:12925
## Max.   :100.0      Max.   :31643      Max.   :21836.0      Max.   :21700
## Room.Board    Books      Personal      PhD
## Min.   :1780      Min.   : 96.0      Min.   : 250      Min.   : 8.00
## 1st Qu.:3597      1st Qu.:470.0      1st Qu.: 850      1st Qu.:62.00
## Median :4200      Median :500.0      Median :1200      Median :75.00
## Mean   :4358      Mean   :549.4      Mean   :1341      Mean   :72.66
## 3rd Qu.:5050      3rd Qu.:600.0      3rd Qu.:1700      3rd Qu.:85.00
## Max.   :8124      Max.   :2340.0      Max.   :6800      Max.   :103.00
## Terminal      S.F.Ratio      perc.alumni      Expend
## Min.   :24.0      Min.   :2.50      Min.   :0.00      Min.   :3186
## 1st Qu.:71.0      1st Qu.:11.50      1st Qu.:13.00      1st Qu.:6751
## Median :82.0      Median :13.60      Median :21.00      Median :8377
## Mean   :79.7      Mean   :14.09      Mean   :22.74      Mean   :9660
## 3rd Qu.:92.0      3rd Qu.:16.50      3rd Qu.:31.00      3rd Qu.:10830
## Max.   :100.0      Max.   :39.80      Max.   :64.00      Max.   :56233
## Grad.Rate      Elite
## Min.   :10.00      No :699
## 1st Qu.:53.00      Yes: 78
## Median :65.00
## Mean   :65.46
## 3rd Qu.:78.00
## Max.   :118.00
```

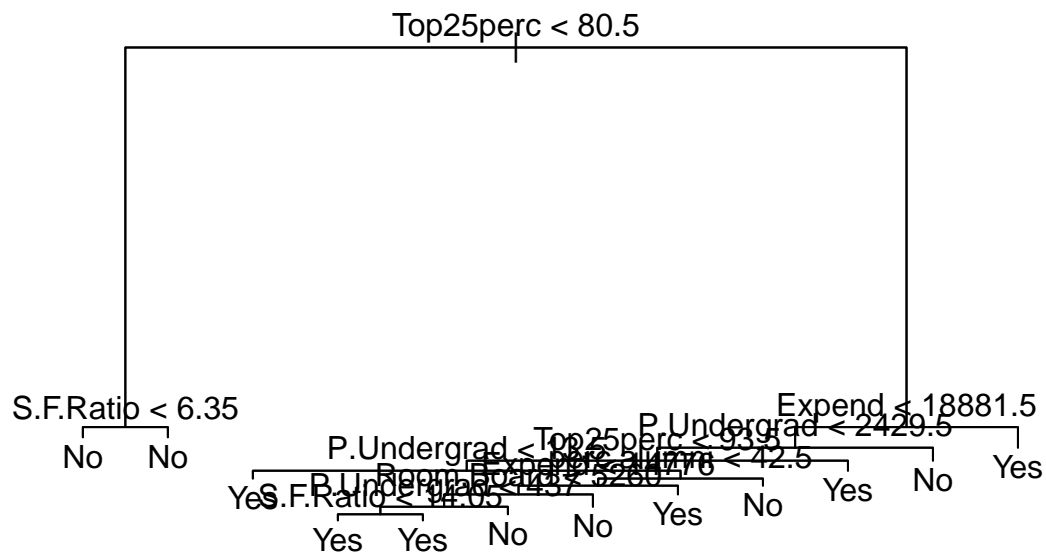
## Q1 a)

First, a decision tree is created, excluding Top10perc.

```
library(tree)
tree_elite = tree(Elite ~ . - Top10perc, data = college)
summary(tree_elite)
```

```
##
## Classification tree:
## tree(formula = Elite ~ . - Top10perc, data = college)
## Variables actually used in tree construction:
## [1] "Top25perc" "S.F.Ratio" "Expend" "P.Undergrad" "perc.alumni"
## [6] "Room.Board"
## Number of terminal nodes: 12
## Residual mean deviance: 0.04263 = 32.61 / 765
## Misclassification error rate: 0.009009 = 7 / 777
```

```
plot(tree_elite, main = "Decision Tree for Elites")
text(tree_elite, pretty = 0)
```



## Q1 b)

Using a set seed for reproducible results, the new tree is trained on 500 random observations from the data set and tested against the remaining observations.

```
set.seed(1)
training_set = sample(1:nrow(college), 500)
test_set = college[-training_set,]
new_tree_elite = tree(Elite ~ . - Top10perc, college, subset = training_set)
summary(new_tree_elite)
```

```
##
## Classification tree:
## tree(formula = Elite ~ . - Top10perc, data = college, subset = training_set)
## Variables actually used in tree construction:
## [1] "Top25perc" "S.F.Ratio" "Expend" "perc.alumni" "Apps"
## [6] "P.Undergrad" "Outstate"
## Number of terminal nodes: 11
## Residual mean deviance: 0.05905 = 28.88 / 489
## Misclassification error rate: 0.014 = 7 / 500
```

```
test_predictions = predict(new_tree_elite, test_set, type = "class")
table = table(test_predictions, test_set$Elite)
print(table)
```

```
##
## test_predictions No Yes
##                No 242  8
##                Yes  7 20
```

```
print ((table[1, 2] + table[2, 1])/sum(table))
```

```
## [1] 0.05415162
```

15 observations are misclassified, giving an error rate of ~5.4%, higher than when the tree was tested on all the data used for fitting - which caused overfitting. This is a more realistic error rate.

### Q1 c)

```
logistic_elite = glm(Elite ~ Top25perc + S.F.Ratio + Expend + P.Undergrad + perc.alumni + Room.Board
+ Outstate + Apps, data = college, family = binomial, subset = training_set)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic_elite)
```

```
##
## Call:
## glm(formula = Elite ~ Top25perc + S.F.Ratio + Expend + P.Undergrad +
##      perc.alumni + Room.Board + Outstate + Apps, family = binomial,
##      data = college, subset = training_set)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.701e+01  6.492e+00  -4.161 3.17e-05 ***
## Top25perc    3.630e-01  7.356e-02   4.935 8.02e-07 ***
## S.F.Ratio   -9.896e-02  1.690e-01  -0.586 0.55813
## Expend       3.584e-04  1.202e-04   2.982 0.00286 **
## P.Undergrad -1.231e-03  7.181e-04  -1.715 0.08639 .
## perc.alumni -1.034e-01  3.806e-02  -2.717 0.00659 **
```

```
## Room.Board -1.044e-03 5.300e-04 -1.970 0.04886 *
## Outstate 8.712e-05 1.423e-04 0.612 0.54027
## Apps 3.793e-05 1.478e-04 0.257 0.79744
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 325.083 on 499 degrees of freedom
## Residual deviance: 61.914 on 491 degrees of freedom
## AIC: 79.914
##
## Number of Fisher Scoring iterations: 10
```

Remove non-significant variables: in this case, 'Apps' has the highest P-value.

```
logistic_elite = glm(Elite ~ Top25perc + S.F.Ratio + Expend + P.Undergrad + perc.alumni + Room.Board
+ Outstate, data = college, family = binomial, subset = training_set)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic_elite)
```

```
##
## Call:
## glm(formula = Elite ~ Top25perc + S.F.Ratio + Expend + P.Undergrad +
##     perc.alumni + Room.Board + Outstate, family = binomial, data = college,
##     subset = training_set)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.773e+01 5.973e+00 -4.643 3.44e-06 ***
## Top25perc    3.686e-01 7.150e-02  5.154 2.54e-07 ***
## S.F.Ratio   -8.484e-02 1.582e-01 -0.536 0.59167
## Expend       3.677e-04 1.142e-04  3.219 0.00129 **
## P.Undergrad -1.111e-03 5.254e-04 -2.115 0.03440 *
## perc.alumni -1.051e-01 3.774e-02 -2.786 0.00533 **
## Room.Board  -1.036e-03 5.298e-04 -1.956 0.05043 .
## Outstate     9.260e-05 1.407e-04  0.658 0.51044
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 325.08 on 499 degrees of freedom
## Residual deviance: 61.98 on 492 degrees of freedom
## AIC: 77.98
##
## Number of Fisher Scoring iterations: 10
```

Remove non-significant variables: in this case, 'S.F.Ratio' has the highest p-value.

```
logistic_elite = glm(Elite ~ Top25perc + Expend + P.Undergrad + perc.alumni + Room.Board
+ Outstate, data = college, family = binomial, subset = training_set)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic_elite)
```

```
##
## Call:
## glm(formula = Elite ~ Top25perc + Expend + P.Undergrad + perc.alumni +
##      Room.Board + Outstate, family = binomial, data = college,
##      subset = training_set)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.919e+01  5.474e+00 -5.333 9.66e-08 ***
## Top25perc    3.662e-01  7.077e-02  5.175 2.28e-07 ***
## Expend       3.810e-04  1.069e-04  3.564 0.000365 ***
## P.Undergrad -1.096e-03  5.198e-04 -2.109 0.034982 *
## perc.alumni -1.033e-01  3.725e-02 -2.772 0.005570 **
## Room.Board  -9.671e-04  5.077e-04 -1.905 0.056815 .
## Outstate     9.697e-05  1.410e-04  0.688 0.491701
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 325.083  on 499  degrees of freedom
## Residual deviance:  62.272  on 493  degrees of freedom
## AIC: 76.272
##
## Number of Fisher Scoring iterations: 10
```

Remove non-significant variables: in this case, 'Outstate' has the highest p-value.

```
logistic_elite = glm(Elite ~ Top25perc + Expend + P.Undergrad + perc.alumni + Room.Board,
data = college, family = binomial, subset = training_set)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic_elite)
```

```
##
## Call:
## glm(formula = Elite ~ Top25perc + Expend + P.Undergrad + perc.alumni +
##      Room.Board, family = binomial, data = college, subset = training_set)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.900e+01  5.480e+00 -5.293 1.21e-07 ***
```

```
## Top25perc      3.619e-01  7.063e-02   5.123 3.00e-07 ***
## Expend         3.993e-04  1.018e-04   3.921 8.83e-05 ***
## P.Undergrad   -1.221e-03  4.988e-04  -2.449 0.01434 *
## perc.alumni   -9.338e-02  3.435e-02  -2.718 0.00656 **
## Room.Board    -7.478e-04  3.770e-04  -1.984 0.04730 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 325.083  on 499  degrees of freedom
## Residual deviance:  62.755  on 494  degrees of freedom
## AIC: 74.755
##
## Number of Fisher Scoring iterations: 10
```

```
predicted_logistic = predict(logistic_elite, newdata = test_set, type = "response")
summary(predicted_logistic)
```

```
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## 0.0000000 0.0000000 0.0000026 0.0875313 0.0009995 0.9999966
```

When using the testing data in the logistic regression, on average, a college has an estimated ~8.75% probability of being elite.

```
predicted_elite = rep("No", 277)
predicted_elite[predicted_logistic > 0.5] = "Yes"
table = table(predicted_elite, test_set$Elite)
print(table)
```

```
##
## predicted_elite  No Yes
##                No 247  5
##                Yes  2 23
```

```
print((table[1, 2] + table[2, 1])/(sum(table)))
```

```
## [1] 0.02527076
```

Using the logistic regression, the error rate is ~2.53%, which is lower than the decision tree and is more accurate in this case.

## Q2

Importing the package and viewing data summary

```
library(ISLR)
View(Auto)
?Auto
```

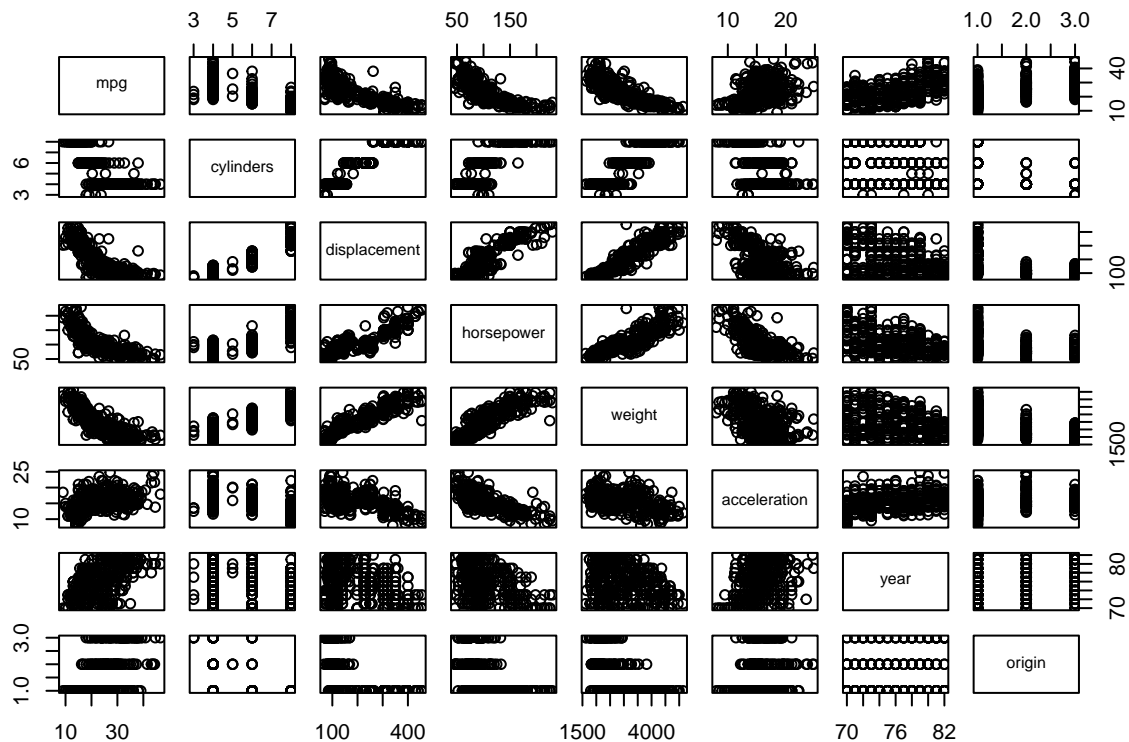
```
## starting httpd help server ... done
```

```
summary(Auto)
```

```
##      mpg      cylinders  displacement  horsepower      weight
## Min.   : 9.00   Min.    :3.000   Min.    : 68.0   Min.    : 46.0   Min.    :1613
## 1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0   1st Qu.:2225
## Median :22.75   Median :4.000   Median :151.0   Median : 93.5   Median :2804
## Mean   :23.45   Mean    :5.472   Mean    :194.4   Mean    :104.5   Mean    :2978
## 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0   3rd Qu.:3615
## Max.   :46.60   Max.    :8.000   Max.    :455.0   Max.    :230.0   Max.    :5140
##
## acceleration      year      origin      name
## Min.    : 8.00   Min.    :70.00   Min.    :1.000   amc matador      : 5
## 1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000   ford pinto       : 5
## Median :15.50   Median :76.00   Median :1.000   toyota corolla   : 5
## Mean    :15.54   Mean     :75.98   Mean    :1.577   amc gremlin      : 4
## 3rd Qu.:17.02   3rd Qu.:79.00   3rd Qu.:2.000   amc hornet       : 4
## Max.    :24.80   Max.     :82.00   Max.    :3.000   chevrolet chevette: 4
##                                     (Other)      :365
```

Q2 a)

```
pairs(Auto[,1:8])
```



## Q2 b)

```
cor(Auto[,1:8])
```

```
##           mpg  cylinders displacement horsepower    weight
## mpg          1.0000000 -0.7776175   -0.8051269 -0.7784268 -0.8322442
## cylinders    -0.7776175  1.0000000    0.9508233  0.8429834  0.8975273
## displacement -0.8051269  0.9508233    1.0000000  0.8972570  0.9329944
## horsepower   -0.7784268  0.8429834    0.8972570  1.0000000  0.8645377
## weight       -0.8322442  0.8975273    0.9329944  0.8645377  1.0000000
## acceleration  0.4233285 -0.5046834   -0.5438005 -0.6891955 -0.4168392
## year         0.5805410 -0.3456474   -0.3698552 -0.4163615 -0.3091199
## origin       0.5652088 -0.5689316   -0.6145351 -0.4551715 -0.5850054
##           acceleration    year    origin
## mpg          0.4233285  0.5805410  0.5652088
## cylinders    -0.5046834 -0.3456474 -0.5689316
## displacement -0.5438005 -0.3698552 -0.6145351
## horsepower   -0.6891955 -0.4163615 -0.4551715
## weight       -0.4168392 -0.3091199 -0.5850054
## acceleration  1.0000000  0.2903161  0.2127458
## year         0.2903161  1.0000000  0.1815277
## origin       0.2127458  0.1815277  1.0000000
```

## Q2 c)

```
mpg_regression = lm(mpg ~ . - name, data = Auto)
summary(mpg_regression)
```

```
##
## Call:
## lm(formula = mpg ~ . - name, data = Auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.5903 -2.1565 -0.1169  1.8690 13.0604
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -17.218435   4.644294  -3.707  0.00024 ***
## cylinders    -0.493376   0.323282  -1.526  0.12780
## displacement  0.019896   0.007515   2.647  0.00844 **
## horsepower   -0.016951   0.013787  -1.230  0.21963
## weight       -0.006474   0.000652  -9.929 < 2e-16 ***
## acceleration  0.080576   0.098845   0.815  0.41548
## year         0.750773   0.050973  14.729 < 2e-16 ***
## origin       1.426141   0.278136   5.127 4.67e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.328 on 384 degrees of freedom
```



```
## Multiple R-squared:  0.8215, Adjusted R-squared:  0.8182
## F-statistic: 252.4 on 7 and 384 DF,  p-value: < 2.2e-16
```

i) Yes, there is a relationship between the predictors and the response. When testing the null hypothesis that all regression coefficients are 0, the F-statistic is returned - with a value of 252.4, suggesting that the overall regression is statistically significant. This is supported by the low p-value.

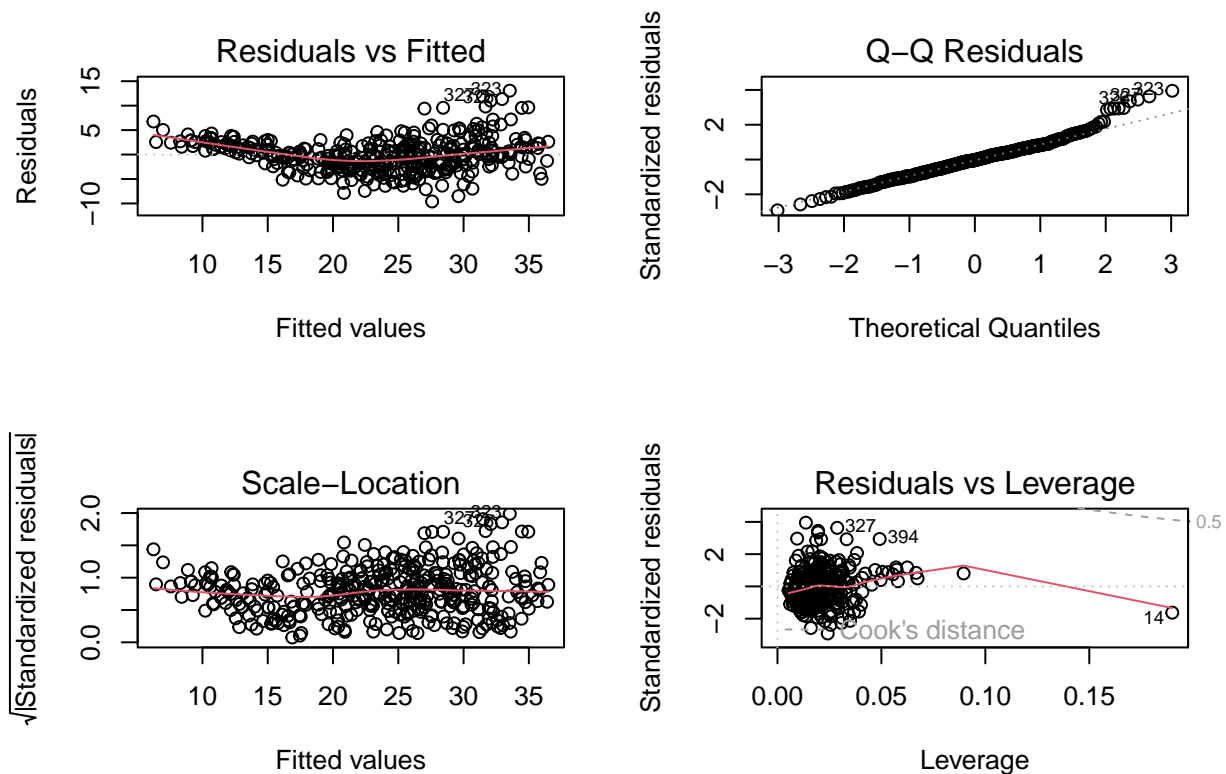
ii) The predictors with a statistically significant relationship are the predictors with low p-values, with  $< 0.05$  being statistically significant. These predictors are: displacement, weight, year and origin. It appears there is a high probability that these regressors affect mpg. These results are supported by the pairs() plot produced earlier, with the exception of horsepower which appears to correlate with mpg in the plot. The lack of statistical significance of acceleration, however, is also supported by the plot.

iii) The coefficients for year and origin are both positive, suggesting a positive relationship between these regressors and mpg. In the case of the year coefficient, an increase in the year results in a  $\sim 0.75$  increase in mpg. Intuitively, this is plausible; later years indicate newer cars and more time for better technology to develop, resulting in efficient cars with higher miles per gallon. The coefficient of origin is  $\sim 1.43$ , almost double the coefficient on year. This suggests the origin of the car has a higher impact on mpg, with American cars being the least efficient, followed by European and Japanese cars.

iv) The insignificant predictors are cylinders, horsepower and acceleration. The first possibility would be to remove the insignificant predictors. This can reduce overfitting and improve generalisability when applying the model to unseen data. This must be done one predictor at a time, in order to observe if previously-insignificant variables become significant. For example, from the plot it can be seen that horsepower and acceleration are highly correlated, so removing one of these predictors can reduce multicollinearity and increase significance, as well as reduce standard error. Another possibility would be to create an interaction term between these two correlated regressors in order to capture the joint effect on mpg, which would avoid omitted variable bias.

Q2 d)

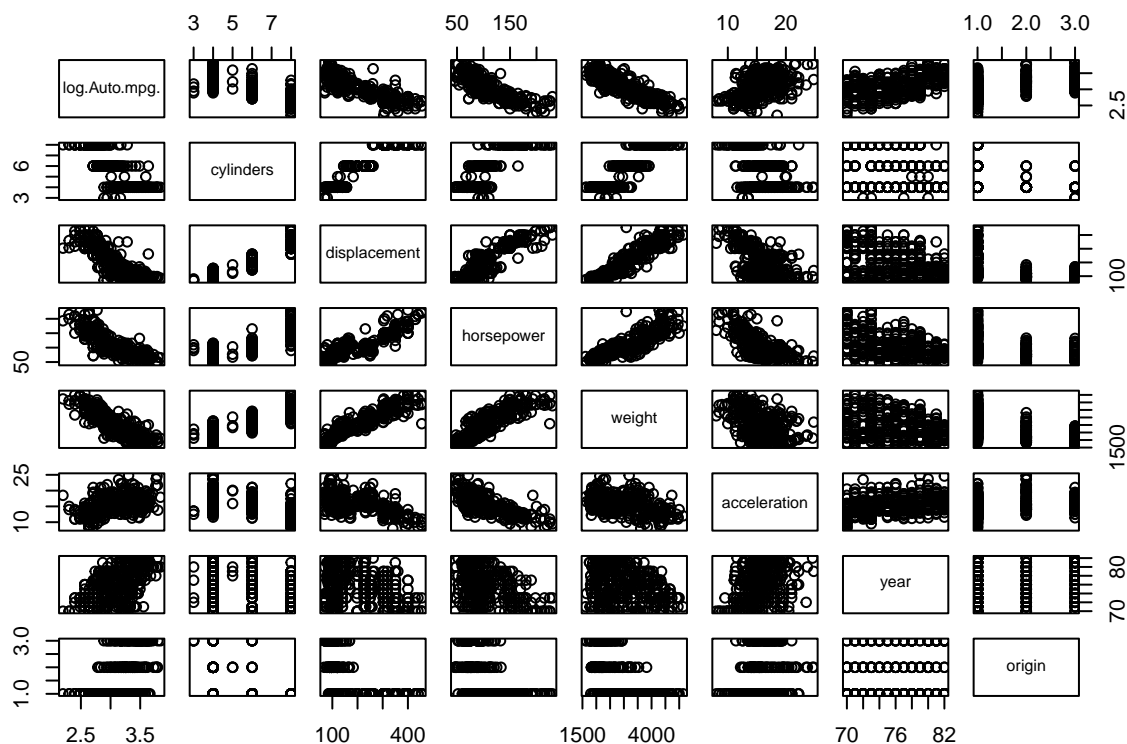
```
par(mfrow=c(2,2))
plot(mpg_regression)
```



From the diagnosis plots, we see the following: - The Residuals vs Fitted plot curves slightly, suggesting that the relationship between the predictors and the response contains some non-linearity. Residuals are higher for smaller fitted values, decreasing as fitted values approach 20, then increasing again. - The Q Q residuals plot shows most residuals have a normal distribution, with the exception of a few outliers towards the end of the graph. - The Scale-Location plot shows some heteroskedasticity, with the same outliers pulling the line upwards due to much higher variance. - The Residuals vs Leverage plot shows clear leverage from point 14, suggesting significant influence on the regression and coefficients.

Q2 e)

```
pairs(data.frame(log(Auto$mpg), Auto[, -c(1,9)]))
```



```
cor(Auto$mpg, Auto$weight)
```

```
## [1] -0.8322442
```

```
cor((log(Auto$mpg)),Auto$weight)
```

```
## [1] -0.8756582
```

Overall, correlations appear to be stronger between log mpg and most variables as seen in the plot. For example, the correlation between log mpg and weight is higher than when mpg is not logged. Log transformations can be more appropriate and result in higher correlation, especially when the relationship was not completely linear to begin with.

Q3 a) i)

```
set.seed(3)
x1=runif(150) # 150 U(0,1) random numbers
x2=0.5*runif(150)+rnorm(150)/5 # rnorm(150) returns 150 N(0,1) random numbers
y=2+2*x1+x2+rnorm(150)
```

$B_0 = 2$ ,  $B_1 = 2$ ,  $B_2 = 1$ ,  $E \sim N(0,1)$

### Q3 a) ii)

```
ytrain = y[1:100]; ytest=y[101:150] # splits y into training and test sets, with 100 and 50 observations
x = data.frame(x1, x2); x.train=x[1:100,]; x.test=x[101:150,]
m1 = lm(ytrain~x1+x2, data=x.train)
summary(m1)
```

```
##
## Call:
## lm(formula = ytrain ~ x1 + x2, data = x.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3324 -0.6809 -0.0203  0.4915  3.5541
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.9065     0.2257   8.449 2.96e-13 ***
## x1            2.0503     0.3761   5.452 3.80e-07 ***
## x2            1.1604     0.4104   2.827 0.00571 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.058 on 97 degrees of freedom
## Multiple R-squared:  0.3108, Adjusted R-squared:  0.2966
## F-statistic: 21.88 on 2 and 97 DF,  p-value: 1.44e-08
```

```
confint.lm(m1)
```

```
##              2.5 %    97.5 %
## (Intercept) 1.4586262 2.354350
## x1          1.3038742 2.796755
## x2          0.3457617 1.975008
```

The coefficient on x1 is 2.0503, while the coefficient on x2 is 1.1604. Both coefficients are significant, with very low p-values. B0 confidence intervals = [1.459, 2.354]. B1 confidence intervals = [1.304, 2.797]. B2 confidence intervals = [0.346, 1.975].

### Q3 b) iii)

```
predict_y = predict(m1, x.test)
summary(predict_y)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1.587   2.751   3.177   3.199   3.535   4.397
```

```
squared_difference = (predict_y - ytest)^2
print (sqrt(mean(squared_difference)))
```

```
## [1] 1.07953
```

The rMSPE is ~1.080.

### Q3 b) i)

```
set.seed(3)
x1=runif(150)
x2=0.5*x1+rnorm(150)/5
y=2+2*x1+x2+rnorm(150)
print (cor(x1, x2))
```

```
## [1] 0.5844431
```

The correlation between x1 and x2 is 0.584.

### Q3 b) ii)

```
ytrain = y[1:100]; ytest=y[101:150] # splits y into training and test sets, with 100 and 50 observations
x = data.frame(x1, x2); x.train=x[1:100,]; x.test=x[101:150,]
m1 = lm(ytrain~x1+x2, data=x.train)
summary(m1)
```

```
##
## Call:
## lm(formula = ytrain ~ x1 + x2, data = x.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3096 -0.6276 -0.0263  0.5668  3.5841
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.3265     0.2030  11.459 < 2e-16 ***
## x1             1.8250     0.4672   3.906 0.000174 ***
## x2             0.4190     0.5197   0.806 0.422167
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.029 on 97 degrees of freedom
## Multiple R-squared:  0.2558, Adjusted R-squared:  0.2405
## F-statistic: 16.67 on 2 and 97 DF,  p-value: 5.981e-07
```

```
confint.lm(m1)
```

```
##              2.5 %    97.5 %
## (Intercept)  1.9235661 2.729523
## x1           0.8976977 2.752343
## x2          -0.6125908 1.450504
```

The coefficient on x1 is 1.8250, while the coefficient on x2 is 0.4190. The coefficient on x2 is not significant, while the x1 coefficient is. B0 confidence intervals = [1.924, 2.730]. B1 confidence intervals = [0.898, 2.752]. B2 confidence intervals = [-0.613, 1.451].

```
predict_y = predict(m1, x.test)
summary(predict_y)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  2.336   2.897   3.332   3.310   3.726   4.201
```

```
squared_difference = (predict_y - ytest)^2
print(sqrt(mean(squared_difference)))
```

```
## [1] 1.090116
```

The rMSPE is 1.090.

### Q3 b) iii)

In the second model run, the coefficient on  $x_2$  is not significant while it was in the first model run.  $x_1$  and  $x_2$  are highly correlated, as  $\text{corr} = 0.584$ . The model is unable to accurately differentiate between the effects of  $x_1$  and  $x_2$  on  $y$ , and as such the coefficient on  $x_2$  becomes statistically insignificant; this is also supported by the larger confidence intervals of the second regression. The new rMSPE, 1.09, is fairly similar to the first regression (rMSPE = 1.07953), suggesting that the difference between actual and predicted values is roughly the same for both regressions.

### Q3 b) iv)

```
m1 = lm(ytrain~x1, data=x.train)
summary(m1)
```

```
##
## Call:
## lm(formula = ytrain ~ x1, data = x.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2686 -0.6139 -0.0734  0.6011  3.5881
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.3182     0.2024  11.453 < 2e-16 ***
## x1            2.0641     0.3604   5.728 1.12e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.027 on 98 degrees of freedom
## Multiple R-squared:  0.2508, Adjusted R-squared:  0.2432
## F-statistic: 32.81 on 1 and 98 DF,  p-value: 1.117e-07
```

The coefficient on  $x_1$  is 2.0641 and is statistically significant.

```
predict_y = predict(m1, x.test)
summary(predict_y)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.403   2.893   3.331   3.319   3.757   4.204
```

```
squared_difference = (predict_y - ytest)^2
print (sqrt(mean(squared_difference)))
```

```
## [1] 1.115781
```

The rMSPE is now 1.116, which is actually slightly higher than the previous regressions. The coefficient on  $x_1$ , 2.0641, is higher than the coefficient on  $x_1$  (1.8250) when the regression with  $x_2$  was also run. This is evidence of multicollinearity, given that the effect of  $x_1$  is higher when  $x_2$  is not included in the regression. The coefficient on  $x_1$  is able to accurately capture the effect of  $x_1$  on  $y$ , without the redundancy of a correlated regressor. Finally, the  $R^2$  is reduced, from 0.2558 in the previous regression to 0.2508. This is expected, as removing a regressor will remove some predictive power of the regression, with a lower proportion of variance being accounted for by the predictors; however, as the reduction is so low (only 0.005) it shows that  $x_2$  is very insignificant in predicting  $y$ .

### Q3 c) i)

```
set.seed(3)
x1=runif(150)
epsilon=rnorm(150)
x2=0.5*runif(150)+epsilon/5
y=2+2*x1+x2+epsilon
cor(x2, epsilon)
```

```
## [1] 0.8217773
```

The correlation between  $x_2$  and the error is 0.822.

### Q3 c) ii)

```
ytrain = y[1:100]; ytest=y[101:150] # splits y into training and test sets, with 100 and 50 observation.
x = data.frame(x1, x2); x.train=x[1:100,]; x.test=x[101:150,]
m1 = lm(ytrain~x1+x2, data=x.train)
summary(m1)
```

```
##
## Call:
## lm(formula = ytrain ~ x1 + x2, data = x.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.35583 -0.43974  0.06475  0.42365  1.19596
```

```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.1972     0.1276   9.383 2.89e-15 ***
## x1            1.8863     0.2103   8.969 2.25e-14 ***
## x2            4.4114     0.2520  17.504 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5911 on 97 degrees of freedom
## Multiple R-squared:  0.8228, Adjusted R-squared:  0.8192
## F-statistic: 225.3 on 2 and 97 DF,  p-value: < 2.2e-16
```

```
confint.lm(m1)
```

```
##           2.5 %    97.5 %
## (Intercept) 0.9439413 1.450410
## x1          1.4689389 2.303755
## x2          3.9111854 4.911557
```

The coefficient on x1 is 1.886, while the coefficient on x2 is 4.411. The coefficients on x1 and x2 are both statistically significant. B0 confidence interval = [0.944, 1.450]. B1 confidence interval = [1.469, 2.304]. B2 confidence interval = [3.911, 4.912].

```
predict_y = predict(m1, x.test)
summary(predict_y)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  1.099   2.455   3.491   3.433   4.476   6.459
```

```
squared_difference = (predict_y - ytest)^2
print (sqrt(mean(squared_difference)))
```

```
## [1] 0.6414308
```

The rMPSE is 0.641.

### Q3) c) iii)

B1 and B2 are both statistically significant. The  $R^2$  is high, at 0.8228, suggesting the predictors capture most of the variance in y. rMSPE is also lower than previously, suggesting the model's estimates are accurate to the true values. This occurs due to endogeneity, causing biased estimates. Additionally, the confidence interval of B2 does not include its true value within the range due to the correlation between B2 and the error term.